

Olu focused on architectural validation and rigorous performance analysis.

This included generating and interpreting the full training/validation accuracy and loss curves across all 15 epochs, checking for any divergence that would indicate overfitting or underfitting, and confirming that the final test-set accuracy of ****99.23%**** (loss 0.0232) is consistent with the validation trajectory. By examining both the quantitative metrics and the shape of the learning curves, Olu ensured the team's interpretation of the results was evidence-based rather than simply accepting a high baseline score. The analysis also verified that each stage of the pipeline—from data loading through evaluation—functioned as intended and that the model learned generalizable digit patterns.